

Plenario

Charlie Catlett, Jonathan Giuffrida, Brett Goldstein, Robert Mitchum*

Abstract

1 Introduction and Landscape

2 Platform

2.1 Automated ETL Builder

2.2 Single Spatio-Temporal Index

2.3 Sensor and Local Data

2.4 API Endpoints

3 Architecture

4 Early Applications

5 Issues Encountered

5.1 Data Issues

5.2 Scaling Issues

5.3 Philosophical Issues

6 Conclusion

Copyright 0000 IEEE. Personal use of this material is permitted. However, permission to reprint/republish this material for advertising or promotional purposes or for creating new collective works for resale or redistribution to servers or lists, or to reuse any copyrighted component of this work in other works must be obtained from the IEEE.

Bulletin of the IEEE Computer Society Technical Committee on Data Engineering

*The Plenario project is funded by the John D. and Catherine T. MacArthur Foundation and the National Science Foundation via an NSF Early-Concept Grant for Exploratory Research (EAGER) for software development (award number 1348865), while the interaction capabilities were driven by the Urban Sciences Research Coordination Network, created with an NSF Building Community and Capacity for Data-Intensive Research in the Social, Behavioral, and Economic Sciences and in Education and Human Resources (BCC-SBE/EHR) award.

All authors: The Urban Center for Computation and Data (<http://urbanccd.org>). Catlett: Argonne National Laboratory (<http://anl.gov>). Catlett and Mitchum: The Computation Institute (<http://ci.uchicago.edu>). Giuffrida and Goldstein: The Harris School for Public Policy Studies, The University of Chicago (<http://harris.uchicago.edu>).

Figure 1: Dataset Lifecycle

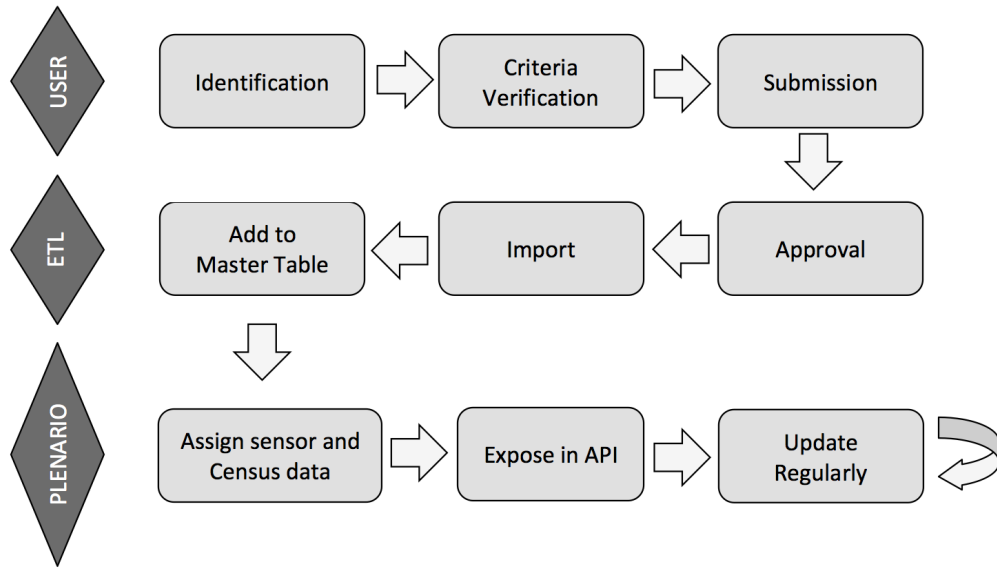


Figure 2: An example search using the user interface

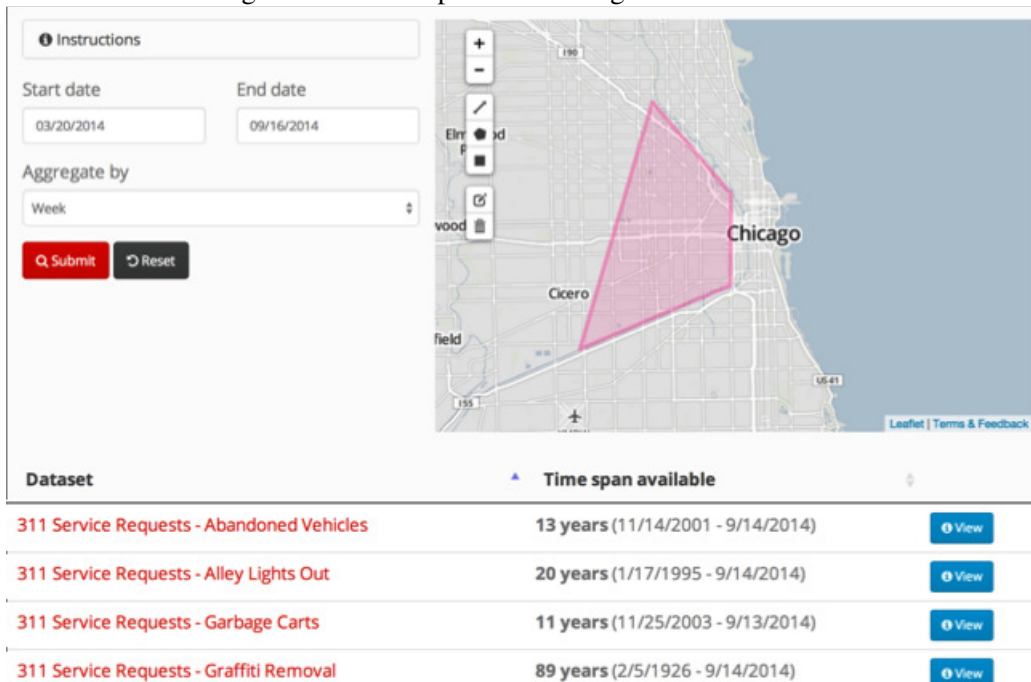


Figure 3: PostgreSQL Schema

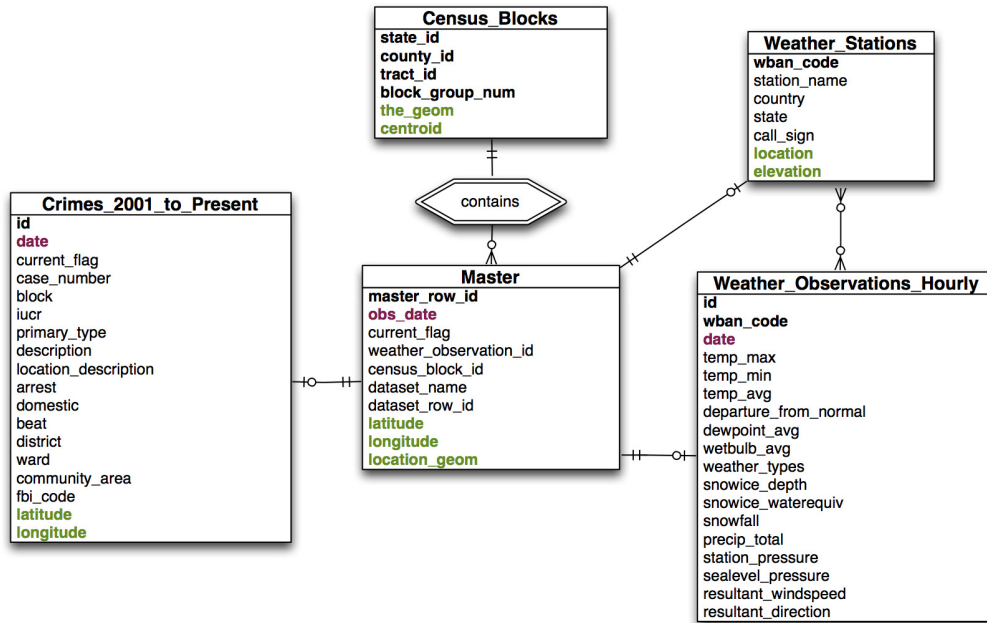


Figure 4: Amazon Web Services Setup

